

Week 5

September 25, 2024

```
[1]: import pandas as pd
df = pd.DataFrame({'Date': ['10/2/2011', '11/2/2011', '12/2/2011', '13/2/2011'],
                  'Event': ['Music', 'Poetry', 'Theatre', 'Comedy'],
                  'Cost': [10000, 5000, 15000, 2000]})
print(df)
```

	Date	Event	Cost
0	10/2/2011	Music	10000
1	11/2/2011	Poetry	5000
2	12/2/2011	Theatre	15000
3	13/2/2011	Comedy	2000

```
[2]: df['Discounted_Price'] = df.apply(lambda row: row.Cost * 0.9, axis = 1)
print(df)
```

	Date	Event	Cost	Discounted_Price
0	10/2/2011	Music	10000	9000.0
1	11/2/2011	Poetry	5000	4500.0
2	12/2/2011	Theatre	15000	13500.0
3	13/2/2011	Comedy	2000	1800.0

```
[3]: df = pd.DataFrame({'Name': ['John', 'Ted', 'Dove', 'Brad', 'Rex'],
                      'Salary': [44000, 35000, 75000, 20000, 6000]})
print(df)
```

	Name	Salary
0	John	44000
1	Ted	35000
2	Dove	75000
3	Brad	20000
4	Rex	6000

```
[4]: def salary_stats(value):
    if value < 10000:
        return "very low"
    elif 10000 <= value < 25000:
        return "low"
    elif 25000 <= value < 40000:
        return "average"
```

```

elif 40000 <= value < 50000:
    return "better"
elif value >= 50000:
    return "very good"

df['salary_stats'] = df['Salary'].map(salary_stats)
df

```

```

[4]:   Name  Salary salary_stats
0  John   44000         better
1   Ted   35000         average
2  Dove   75000        very good
3  Brad   20000           low
4   Rex    6000        very low

```

```

[9]: import pandas as pd
data = pd.DataFrame({
    'Name' : ['A', 'B', 'C', 'D', 'E', 'F'],
    'Education' : ['High School', 'Masters', 'Doctorate', 'Bachelors', 'Masters', 'High School']})
data

```

```

[9]:   Name  Education
0    A  High School
1    B    Masters
2    C  Doctorate
3    D  Bachelors
4    E    Masters
5    F  High School

```

1 Binary Encoding

```

[10]: education_data = pd.get_dummies(data.Education)
print(education_data)

```

```

   Bachelors  Doctorate  High School  Masters
0     False     False         True    False
1     False     False         False     True
2     False      True         False    False
3      True     False         False    False
4     False     False         False     True
5     False     False         True    False

```

2 Ranking Transformation

```
[11]: education_map = {
      'High School' : 1,
      'Bachelors' : 2,
      'Masters': 3,
      'Doctorate': 4
      }
      education_data = data['Education'].map(education_map)
      data['Education'] = education_data
      data
```

```
[11]:   Name  Education
      0    A         1
      1    B         3
      2    C         4
      3    D         2
      4    E         3
      5    F         1
```

```
[12]: education_map = {
      'High School' : 12,
      'Bachelors' : 16,
      'Masters': 18,
      'Doctorate': 21
      }
      education_data = data['Education'].map(education_map)
      data['Education'] = education_data
      data
```

```
[12]:   Name  Education
      0    A        NaN
      1    B        NaN
      2    C        NaN
      3    D        NaN
      4    E        NaN
      5    F        NaN
```

3 Adding data objects- rows

```
[13]: df.loc[len(df.index)]=['Hruthvik', 15000, 'low']
      df
```

```
[13]:   Name  Salary salary_stats
      0   John   44000      better
      1    Ted   35000    average
      2   Dove   75000  very good
```

3	Brad	20000	low
4	Rex	6000	very low
5	Hruthvik	15000	low

4 Combining two data frames

```
[15]: import pandas as pd
d1 = {'Name': ['Pankaj', 'Meghna', 'Lisa'], 'Country': ['India', 'India', 'USA']}
df1 = pd.DataFrame(d1)
print('DataFrame 1:\n', df1, '\n')
df2 = pd.DataFrame({'ID': [1, 2, 3], 'Name': ['Pankaj', 'Anupam', 'Amit']})
print('DataFrame 2:\n', df2, '\n')
df3 = pd.DataFrame({'Name': ['Priya'], 'Country': ['India'], 'Role': ['COO']})
print('DataFrame 3:\n', df3, '\n')
```

```
DataFrame 1:
      Name Country
0  Pankaj   India
1  Meghna   India
2    Lisa    USA
```

```
DataFrame 2:
   ID  Name
0   1  Pankaj
1   2  Anupam
2   3   Amit
```

```
DataFrame 3:
      Name Country Role
0  Priya   India   COO
```

```
[16]: same_cols_df = pd.concat([df1, df3], ignore_index=True)
same_cols_df
```

```
[16]:      Name Country Role
0  Pankaj   India  NaN
1  Meghna   India  NaN
2    Lisa    USA  NaN
3  Priya   India  COO
```

```
[17]: a_df=df1.append(df2, ignore_index=True)
a_df
```

AttributeError

Traceback (most recent call last)

```

~\AppData\Local\Temp\ipykernel_26088\2048347772.py in ?()
----> 1 a_df=df1.append(df2, ignore_index=True)
      2 a_df

~\anaconda3\Lib\site-packages\pandas\core\generic.py in ?(self, name)
    6295         and name not in self._accessors
    6296         and self._info_axis.
↪ _can_hold_identifiers_and_holds_name(name)
    6297     ):
    6298         return self[name]
-> 6299     return object.__getattr__(self, name)

AttributeError: 'DataFrame' object has no attribute 'append'

```

```

[18]: c_df = pd.concat([df1,df2],ignore_index=True)
      c_df

```

```

[18]:
   Name Country  ID
0  Pankaj   India NaN
1  Meghna   India NaN
2    Lisa    USA  NaN
3  Pankaj    NaN  1.0
4  Anupam    NaN  2.0
5    Amit    NaN  3.0

```

5 Defaut Mergeing - inner join

```

[19]: df_merged = df1.merge(df2)
      print('Result:\n', df_merged)

```

```

Result:
   Name Country  ID
0  Pankaj   India  1

```

6 Mergeing DataFrames with left, Right and outer join

```

[20]: print('Result Left Join:\n', df1.merge(df2, how='left'))
      print('Result Right Join:\n', df1.merge(df2, how='right'))
      print('Result Outer Join:\n', df1.merge(df2, how='outer'))

```

```

Result Left Join:
   Name Country  ID
0  Pankaj   India  1.0
1  Meghna   India  NaN
2    Lisa    USA  NaN
Result Right Join:

```

	Name	Country	ID
0	Pankaj	India	1
1	Anupam	NaN	2
2	Amit	NaN	3

Result Outer Join:

	Name	Country	ID
0	Amit	NaN	3.0
1	Anupam	NaN	2.0
2	Lisa	USA	NaN
3	Meghna	India	NaN
4	Pankaj	India	1.0

7 Merging dataframes with specific columns

```
[30]: import pandas as pd

dict1 = {
    'ID': [1, 2, 3],
    'Name': ['Pankaj', 'Meghna', 'Lisa'],
    'Country': ['India', 'India', 'India'],
    'Role': ['CEO', 'CTO', 'CTO']
}

df1 = pd.DataFrame(dict1)
df2 = pd.DataFrame({'ID': [1, 2, 3], 'Name': ['Pankaj', 'Anupam', 'Amit']})

print(df1.merge(df2, on='ID'))
print('\n', df1.merge(df2, on='Name'))
```

	ID	Name_x	Country	Role	Name_y
0	1	Pankaj	India	CEO	Pankaj
1	2	Meghna	India	CTO	Anupam
2	3	Lisa	India	CTO	Amit

	ID_x	Name	Country	Role	ID_y
0	1	Pankaj	India	CEO	1

8 Titanic CSV

```
[41]: import pandas as pd
import missingno as msno
%matplotlib inline
```

```
[34]: titanic_df = pd.read_csv("titanic.csv")
titanic_df
```

```
[34]: PassengerId  Survived  Pclass  \
0            1         0         3
1            2         1         1
2            3         1         3
3            4         1         1
4            5         0         3
..         ...         ...         ...
886         887         0         2
887         888         1         1
888         889         0         3
889         890         1         1
890         891         0         3
```

```

                                Name      Sex  Age  SibSp  \
0                        Braund, Mr. Owen Harris    male  22.0      1
1  Cumings, Mrs. John Bradley (Florence Briggs Th... female  38.0      1
2                        Heikkinen, Miss. Laina    female  26.0      0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)    female  35.0      1
4                        Allen, Mr. William Henry    male  35.0      0
..         ...         ...         ...         ...
886                        Montvila, Rev. Juozas    male  27.0      0
887                        Graham, Miss. Margaret Edith    female  19.0      0
888      Johnston, Miss. Catherine Helen "Carrie"    female   NaN      1
889                        Behr, Mr. Karl Howell    male  26.0      0
890                        Dooley, Mr. Patrick    male  32.0      0
```

```

      Parch      Ticket    Fare Cabin Embarked
0         0      A/5 21171    7.2500   NaN      S
1         0         PC 17599   71.2833   C85      C
2         0  STON/O2. 3101282    7.9250   NaN      S
3         0      113803   53.1000  C123      S
4         0      373450    8.0500   NaN      S
..         ...         ...         ...         ...
886         0      211536   13.0000   NaN      S
887         0      112053   30.0000   B42      S
888         2      W./C. 6607   23.4500   NaN      S
889         0      111369   30.0000  C148      C
890         0      370376    7.7500   NaN      Q
```

[891 rows x 12 columns]

```
[35]: titanic_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
#   Column      Non-Null Count  Dtype
```

```

0  PassengerId  891 non-null    int64
1  Survived     891 non-null    int64
2  Pclass       891 non-null    int64
3  Name         891 non-null    object
4  Sex          891 non-null    object
5  Age          714 non-null    float64
6  SibSp        891 non-null    int64
7  Parch        891 non-null    int64
8  Ticket       891 non-null    object
9  Fare         891 non-null    float64
10 Cabin        204 non-null    object
11 Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB

```

```
[36]: titanic_df.isnull()
```

```

[36]:      PassengerId  Survived  Pclass   Name   Sex    Age  SibSp  Parch  Ticket \
0             False     False   False  False  False  False  False  False  False
1             False     False   False  False  False  False  False  False  False
2             False     False   False  False  False  False  False  False  False
3             False     False   False  False  False  False  False  False  False
4             False     False   False  False  False  False  False  False  False
..            ...      ...    ...    ...    ...    ...    ...    ...
886           False     False   False  False  False  False  False  False  False
887           False     False   False  False  False  False  False  False  False
888           False     False   False  False  False   True  False  False  False
889           False     False   False  False  False  False  False  False  False
890           False     False   False  False  False  False  False  False  False


      Fare  Cabin  Embarked
0     False   True     False
1     False  False     False
2     False   True     False
3     False  False     False
4     False   True     False
..      ...    ...      ...
886  False   True     False
887  False  False     False
888  False   True     False
889  False  False     False
890  False   True     False

[891 rows x 12 columns]

```


```
[42]: msno.bar(titanic_df)
```

```
[42]: <Axes: >
```

```
[43]: msno.matrix(titanic_df)
```

```
[43]: <Axes: >
```



```
[44]: titanic_df.isnull().sum()
```

```
[44]: PassengerId    0
      Survived      0
      Pclass       0
      Name         0
      Sex          0
      Age         177
```

```
SibSp      0
Parch      0
Ticket     0
Fare       0
Cabin     687
Embarked    2
dtype: int64
```

```
[45]: df = titanic_df.dropna(axis=0)
df.isnull().sum()
```

```
[45]: PassengerId    0
Survived          0
Pclass            0
Name              0
Sex               0
Age              0
SibSp            0
Parch            0
Ticket           0
Fare             0
Cabin            0
Embarked          0
dtype: int64
```

```
[46]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 183 entries, 1 to 889
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     183 non-null   int64
1   Survived        183 non-null   int64
2   Pclass          183 non-null   int64
3   Name            183 non-null   object
4   Sex             183 non-null   object
5   Age             183 non-null   float64
6   SibSp           183 non-null   int64
7   Parch           183 non-null   int64
8   Ticket          183 non-null   object
9   Fare            183 non-null   float64
10  Cabin           183 non-null   object
11  Embarked        183 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 18.6+ KB
```

```
[47]: titanic_df.columns
```

```
[47]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
          'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],
          dtype='object')
```

```
[50]: df = titanic_df.drop(['Pclass'],axis=1)
      df.isnull().sum()
```

```
[50]: PassengerId      0
      Survived        0
      Name           0
      Sex            0
      Age           177
      SibSp          0
      Parch          0
      Ticket         0
      Fare           0
      Cabin         687
      Embarked       2
      dtype: int64
```

```
[51]: titanic_df['Pclass'].unique()
```

```
[51]: array([3, 1, 2], dtype=int64)
```

```
[53]: titanic_df['Pclass'] = titanic_df['Pclass'].fillna('C')
```

```
[54]: titanic_df['Pclass'].isnull().sum()
```

```
[54]: 0
```

```
[55]: titanic_df
```

```
[55]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	
..	
886	887	0	2	
887	888	1	1	
888	889	0	3	
889	890	1	1	
890	891	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	

2		Heikkinen, Miss. Laina	female	26.0	0
3		Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1
4		Allen, Mr. William Henry	male	35.0	0
..	
886		Montvila, Rev. Juozas	male	27.0	0
887		Graham, Miss. Margaret Edith	female	19.0	0
888		Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1
889		Behr, Mr. Karl Howell	male	26.0	0
890		Dooley, Mr. Patrick	male	32.0	0

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S
..
886	0	211536	13.0000	NaN	S
887	0	112053	30.0000	B42	S
888	2	W./C. 6607	23.4500	NaN	S
889	0	111369	30.0000	C148	C
890	0	370376	7.7500	NaN	Q

[891 rows x 12 columns]

```
[57]: mean = titanic_df['Age'].mean()
print(mean)
#Replace the missing values for numerical columns with mean
titanic_df['Age'] = titanic_df['Age'].fillna(mean)
titanic_df['Age']
```

29.69911764705882

```
[57]: 0      22.000000
      1      38.000000
      2      26.000000
      3      35.000000
      4      35.000000
      ...
      886     27.000000
      887     19.000000
      888     29.699118
      889     26.000000
      890     32.000000
      Name: Age, Length: 891, dtype: float64
```

```
[58]: titanic_df = pd.read_csv("titanic.csv")
      #Replace the missing values for categorical columns with mode
      mode = titanic_df['Pclass'].mode()[0]
      print(mode)
      titanic_df['Pclass'] = titanic_df['Pclass'].fillna(mode)
```

3

```
[59]: titanic_df['Pclass']
```

```
[59]: 0      3
      1      1
      2      3
      3      1
      4      3
      ..
      886    2
      887    1
      888    3
      889    1
      890    3
      Name: Pclass, Length: 891, dtype: int64
```

```
[61]: titanic_df['Age'] = titanic_df['Age'].fillna(titanic_df['Age'].median())
      titanic_df['Age']
```

```
[61]: 0      22.0
      1      38.0
      2      26.0
      3      35.0
      4      35.0
      ...
      886    27.0
      887    19.0
      888    28.0
      889    26.0
      890    32.0
      Name: Age, Length: 891, dtype: float64
```

```
[62]: titanic_df = pd.read_csv("titanic.csv")
      titanic_df
```

```
[62]:   PassengerId  Survived  Pclass  \
0             1         0        3
1             2         1        1
2             3         1        3
3             4         1        1
4             5         0        3
```

```

..      ...      ...      ...
886      887      0      2
887      888      1      1
888      889      0      3
889      890      1      1
890      891      0      3

```

```

                                Name      Sex  Age  SibSp  \
0      Braund, Mr. Owen Harris      male  22.0      1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0      1
2      Heikkinen, Miss. Laina      female  26.0      0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)  female  35.0      1
4      Allen, Mr. William Henry      male  35.0      0
..
886      Montvila, Rev. Juozas      male  27.0      0
887      Graham, Miss. Margaret Edith      female  19.0      0
888  Johnston, Miss. Catherine Helen "Carrie"      female   NaN      1
889      Behr, Mr. Karl Howell      male  26.0      0
890      Dooley, Mr. Patrick      male  32.0      0

```

```

      Parch      Ticket      Fare Cabin Embarked
0         0      A/5 21171      7.2500   NaN      S
1         0      PC 17599     71.2833   C85      C
2         0  STON/O2. 3101282      7.9250   NaN      S
3         0      113803     53.1000  C123      S
4         0      373450      8.0500   NaN      S
..
886         0      211536     13.0000   NaN      S
887         0      112053     30.0000   B42      S
888         2      W./C. 6607     23.4500   NaN      S
889         0      111369     30.0000  C148      C
890         0      370376      7.7500   NaN      Q

```

[891 rows x 12 columns]

```
[63]: new_df = titanic_df.fillna(method="ffill")
      new_df
```

```

C:\Users\skand\AppData\Local\Temp\ipykernel_26088\3071186871.py:1:
FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a
future version. Use obj.ffill() or obj.bfill() instead.
      new_df = titanic_df.fillna(method="ffill")

```

```
[63]: PassengerId  Survived  Pclass  \
0             1         0         3
1             2         1         1
2             3         1         3
3             4         1         1

```

4	5	0	3
..
886	887	0	2
887	888	1	1
888	889	0	3
889	890	1	1
890	891	0	3

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	
..	
886	Montvila, Rev. Juozas	male	27.0	0	
887	Graham, Miss. Margaret Edith	female	19.0	0	
888	Johnston, Miss. Catherine Helen "Carrie"	female	19.0	1	
889	Behr, Mr. Karl Howell	male	26.0	0	
890	Dooley, Mr. Patrick	male	32.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	C85	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	C123	S
..	
886	0	211536	13.0000	C50	S
887	0	112053	30.0000	B42	S
888	2	W./C. 6607	23.4500	B42	S
889	0	111369	30.0000	C148	C
890	0	370376	7.7500	C148	Q

[891 rows x 12 columns]

```
[64]: new_df = titanic_df.fillna(method="ffill",limit=1)
new_df
```

```
C:\Users\skand\AppData\Local\Temp\ipykernel_26088\3418266207.py:1:
FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a
future version. Use obj.ffill() or obj.bfill() instead.
new_df = titanic_df.fillna(method="ffill",limit=1)
```

```
[64]: PassengerId  Survived  Pclass  \
0             1         0         3
1             2         1         1
2             3         1         3
```

3	4	1	1
4	5	0	3
..
886	887	0	2
887	888	1	1
888	889	0	3
889	890	1	1
890	891	0	3

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	
..	
886	Montvila, Rev. Juozas	male	27.0	0	
887	Graham, Miss. Margaret Edith	female	19.0	0	
888	Johnston, Miss. Catherine Helen "Carrie"	female	19.0	1	
889	Behr, Mr. Karl Howell	male	26.0	0	
890	Dooley, Mr. Patrick	male	32.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	C85	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	C123	S
..
886	0	211536	13.0000	NaN	S
887	0	112053	30.0000	B42	S
888	2	W./C. 6607	23.4500	B42	S
889	0	111369	30.0000	C148	C
890	0	370376	7.7500	C148	Q

[891 rows x 12 columns]

```
[65]: new_df = titanic_df.fillna(method="bfill")
new_df
```

C:\Users\skand\AppData\Local\Temp\ipykernel_26088\348161704.py:1: FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.

```
new_df = titanic_df.fillna(method="bfill")
```

```
[65]: PassengerId  Survived  Pclass  \
0             1           0         3
1             2           1         1
```


2	3	1	3
3	4	1	1
4	5	0	3
..
886	887	0	2
887	888	1	1
888	889	0	3
889	890	1	1
890	891	0	3

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	
..	
886	Montvila, Rev. Juozas	male	27.0	0	
887	Graham, Miss. Margaret Edith	female	19.0	0	
888	Johnston, Miss. Catherine Helen "Carrie"	female	26.0	1	
889	Behr, Mr. Karl Howell	male	26.0	0	
890	Dooley, Mr. Patrick	male	32.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	C85	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	C123	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	E46	S
..
886	0	211536	13.0000	B42	S
887	0	112053	30.0000	B42	S
888	2	W./C. 6607	23.4500	C148	S
889	0	111369	30.0000	C148	C
890	0	370376	7.7500	NaN	Q

[891 rows x 12 columns]

9 Numerosity Data Reduction

```
[66]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```

```
[67]: customer_df = pd.read_csv('customer_churn.csv')
print(customer_df.shape)
print(customer_df.Churn.value_counts())
```

```
(900, 10)
Churn
0    750
1    150
Name: count, dtype: int64
```

```
[69]: import pandas as pd

# Example DataFrame creation (replace with your actual DataFrame)
customer_df = pd.DataFrame({
    'Churn': [0, 1] * 500, # Example data
    'Feature1': range(1000), # Example feature
    'Feature2': range(1000) # Example feature
})

# Ensure the sample size is appropriate
sample_size = min(1000, customer_df.shape[0])

# Sample the DataFrame
customer_df_rs = customer_df.sample(sample_size, random_state=1)

# Separate features and target
y = customer_df_rs['Churn']
Xs = customer_df_rs.drop(columns=['Churn'])

print(customer_df_rs.shape)
```

```
(1000, 3)
```

```
[70]: customer_df_rs
```

```
[70]:
```

	Churn	Feature1	Feature2
507	1	507	507
818	0	818	818
452	0	452	452
368	0	368	368
242	0	242	242
..
767	1	767	767
72	0	72	72
908	0	908	908
235	1	235	235
37	1	37	37

[1000 rows x 3 columns]

```
[71]: print(customer_df_rs.Churn.value_counts())
```

```
Churn
1     500
0     500
Name: count, dtype: int64
```

10 Stratified Sampling


```
[72]: n,s=len(customer_df),1000
      print(n,s)
      r = s/n
      print('Ratio of each Churn class in sample:',r)
      sample_df = customer_df.groupby('Churn').apply(lambda sdf: sdf.
        ↳sample(round(len(sdf))))
      print(sample_df.Churn.value_counts())
```

```
1000 1000
Ratio of each Churn class in sample: 1.0
Churn
0     500
1     500
Name: count, dtype: int64
```

```
C:\Users\skand\AppData\Local\Temp\ipykernel_26088\3439575783.py:5:
DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns.
This behavior is deprecated, and in a future version of pandas the grouping
columns will be excluded from the operation. Either pass `include_groups=False`
to exclude the groupings or explicitly select the grouping columns after groupby
to silence this warning.
      sample_df = customer_df.groupby('Churn').apply(lambda sdf:
sdf.sample(round(len(sdf))))
```


```
[73]: customer_df.Churn.value_counts().plot.bar()
```

```
[73]: <Axes: xlabel='Churn'>
```



```
[74]: sample_df.Churn.value_counts().plot.bar()
```

```
[74]: <Axes: xlabel='Churn'>
```



11 Random Over/Under sampling


```
[75]: n,s=len(customer_df),500
sample_df = customer_df.groupby('Churn').apply(lambda sdf: sdf.sample(250))
print(sample_df.Churn.value_counts())
```

```
Churn
0    250
1    250
Name: count, dtype: int64
```

```
C:\Users\skand\AppData\Local\Temp\ipykernel_26088\3198813423.py:2:
DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns.
This behavior is deprecated, and in a future version of pandas the grouping
columns will be excluded from the operation. Either pass `include_groups=False`
to exclude the groupings or explicitly select the grouping columns after groupby
to silence this warning.
sample_df = customer_df.groupby('Churn').apply(lambda sdf: sdf.sample(250))
```

```
[76]: sample_df.Churn.value_counts().plot.bar()
```

```
[76]: <Axes: xlabel='Churn'>
```



```
[77]: sample_df
```

```
[77]:
```

	Churn	Feature1	Feature2
Churn			
0	136	0	136
	134	0	134
	892	0	892
	784	0	784
	970	0	970
...
1	383	1	383
	305	1	305
	559	1	559
	115	1	115
	11	1	11

```
[500 rows x 3 columns]
```

12 Outliers Detection

```
[78]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.cm as cm
```

```
[79]: titanic_df = pd.read_csv("titanic.csv")
titanic_df
```

```
[79]: PassengerId  Survived  Pclass  \
0                1         0       3
1                2         1       1
2                3         1       3
3                4         1       1
4                5         0       3
..            ...     ...     ...
886            887         0       2
887            888         1       1
888            889         0       3
889            890         1       1
890            891         0       3
```

```

                                Name      Sex  Age  SibSp  \
0                Braund, Mr. Owen Harris   male  22.0     1
1  Cumings, Mrs. John Bradley (Florence Briggs Th... female  38.0     1
2                Heikkinen, Miss. Laina   female  26.0     0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)   female  35.0     1
4                Allen, Mr. William Henry   male  35.0     0
..            ...     ...     ...     ...
886                Montvila, Rev. Juozas   male  27.0     0
887                Graham, Miss. Margaret Edith   female  19.0     0
888    Johnston, Miss. Catherine Helen "Carrie"   female   NaN     1
889                Behr, Mr. Karl Howell   male  26.0     0
890                Dooley, Mr. Patrick   male  32.0     0
```

```


    Parch      Ticket    Fare Cabin Embarked
0      0      A/5 21171    7.2500   NaN      S
1      0      PC 17599   71.2833   C85      C
2      0  STON/O2. 3101282    7.9250   NaN      S
3      0      113803   53.1000  C123      S
4      0      373450    8.0500   NaN      S
..     ...     ...     ...     ...
886     0      211536   13.0000   NaN      S
887     0      112053   30.0000   B42      S
888     2      W./C. 6607   23.4500   NaN      S
889     0      111369   30.0000  C148      C
```

890 0 370376 7.7500 NaN Q

[891 rows x 12 columns]

13 Scatter plot to detect outliers


```
[81]: fig,ax = plt.subplots(figsize=(10,4))
      ax.scatter(titanic_df['Age'],titanic_df['Fare'])
      ax.set_xlabel('Age')
      ax.set_ylabel('Fare')
      plt.title("Scatter plot")
      plt.show()
```



14 Box plot to detect outliers

```
[83]: titanic_df['Age'].plot(kind='box')
```

```
[83]: <Axes: >
```

```
[85]: q1 = titanic_df["Age"].quantile(0.25)
      # finding the 3rd quartile
      q3 = titanic_df['Age'].quantile(0.75)
      # finding the iqr region
      iqr = q3-q1
      # finding upper and lower whiskers
      upper_bound = q3+(1.5*iqr)
      lower_bound = q1-(1.5*iqr)
```

```
[86]: age_arr = titanic_df["Age"]
      outliers = age_arr[(age_arr <= lower_bound) | (age_arr >= upper_bound)]
      print('The following are the outliers in the boxplot of age:\n',outliers)
```

The following are the outliers in the boxplot of age:


33	66.0
54	65.0
96	71.0
116	70.5
280	65.0
456	65.0
493	71.0
630	80.0

```
672    70.0
745    70.0
851    74.0
Name: Age, dtype: float64
```

15 Histogram to detect outliers

```
[88]: titanic_df['Fare'].plot(kind='hist')
```

```
[88]: <Axes: ylabel='Frequency'>
```



16 Remove Data Objects with outliers

```
[90]: upperIndex = titanic_df[titanic_df['Age']>upper_bound].index
titanic_df.drop(upperIndex,inplace=True)
lowerIndex = titanic_df[titanic_df['Age']<lower_bound].index
titanic_df.drop(lowerIndex,inplace=True)
titanic_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 880 entries, 0 to 890
```

Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	PassengerId	880 non-null	int64
1	Survived	880 non-null	int64
2	Pclass	880 non-null	int64
3	Name	880 non-null	object
4	Sex	880 non-null	object
5	Age	703 non-null	float64
6	SibSp	880 non-null	int64
7	Parch	880 non-null	int64
8	Ticket	880 non-null	object
9	Fare	880 non-null	float64
10	Cabin	199 non-null	object
11	Embarked	878 non-null	object

dtypes: float64(2), int64(5), object(5)

memory usage: 89.4+ KB

17 Replcing Outliners with upper/lower caps

```
[91]: titanic_df = pd.read_csv("titanic.csv")
```

```
[93]: fare_arr = titanic_df["Fare"]
upper_cap = np.percentile(fare_arr,1)
lower_cap = np.percentile(fare_arr,99)
outliers = fare_arr[(fare_arr < upper_cap) | (fare_arr > lower_cap)]
print('The following are the outliers in the boxplot of fare:\n',outliers)
```

The following are the outliers in the boxplot of fare:

```
27      263.0000
88      263.0000
258     512.3292
311     262.3750
341     263.0000
438     263.0000
679     512.3292
737     512.3292
742     262.3750
```

Name: Fare, dtype: float64

```
[96]: for i in titanic_df['Fare']:
        if i<lower_bound :
            titanic_df['Fare'] = titanic_df['Fare'].replace(i,lower_cap)
        elif i>upper_bound :
            titanic_df['Fare'] = titanic_df['Fare'].replace(i,upper_cap)
```

```
[97]: titanic_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId      891 non-null    int64
1   Survived         891 non-null    int64
2   Pclass           891 non-null    int64
3   Name             891 non-null    object
4   Sex              891 non-null    object
5   Age              714 non-null    float64
6   SibSp            891 non-null    int64
7   Parch           891 non-null    int64
8   Ticket           891 non-null    object
9   Fare             891 non-null    float64
10  Cabin            204 non-null    object
11  Embarked         889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

18 replacing outliers with mean

```
[98]: titanic_df = pd.read_csv("titanic.csv")
```

```
[101]: m = np.mean(titanic_df['Age'])
print('mean:',m)
for i in titanic_df['Age']:
    if i<lower_bound or i>upper_bound :
        titanic_df['Age'] = titanic_df['Age'].replace(i,m)
```

```
mean: 29.69911764705882
```

19 Replacing Outliers with median

```
[2]: import pandas as pd
titanic_df = pd.read_csv("titanic.csv")
```

```
[5]: q1 = titanic_df["Age"].quantile(0.25)
     # finding the 3rd quartile
q3 = titanic_df['Age'].quantile(0.75)
     # finding the iqr region
iqr = q3-q1
     # finding upper and lower whiskers
upper_bound = q3+(1.5*iqr)
lower_bound = q1-(1.5*iqr)
```

```
[8]: m = titanic_df['Age'].median()
      print(m)
      for i in titanic_df['Age']:
          if i<lower_bound or i>upper_bound :
              titanic_df['Age'] = titanic_df['Age'].replace(i,m)
```

28.0

```
[ ]:
```