

Week-10

September 26, 2024

```
[2]: # Import necessary libraries
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.datasets import mnist
import matplotlib.pyplot as plt

# Load the MNIST dataset
(x_train, y_train), (x_test, y_test) = mnist.load_data()

# Preprocessing: Normalize the images to [0, 1] range and reshape them
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0
x_train = x_train.reshape(-1, 28, 28, 1)
x_test = x_test.reshape(-1, 28, 28, 1)

# Convert class labels to one-hot encoded vectors
y_train = tf.keras.utils.to_categorical(y_train, 10)
y_test = tf.keras.utils.to_categorical(y_test, 10)

# Build the neural network model
model = models.Sequential()

# First convolutional layer
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2)))

# Second convolutional layer
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))

# Third convolutional layer
model.add(layers.Conv2D(64, (3, 3), activation='relu'))

# Flatten the output and add a fully connected layer
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
```

```

# Output layer with 10 units (one for each digit)
model.add(layers.Dense(10, activation='softmax'))

# Compile the model
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(x_train, y_train, epochs=5, batch_size=64,
                    validation_split=0.1)

# Evaluate the model on the test dataset
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
print(f'Test accuracy: {test_acc:.4f}')

# Plot training & validation accuracy values
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# Plot training & validation loss values
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()


```


```

Epoch 1/5
844/844          34s 35ms/step -
accuracy: 0.8613 - loss: 0.4561 - val_accuracy: 0.9813 - val_loss: 0.0598
Epoch 2/5
844/844          29s 35ms/step -
accuracy: 0.9829 - loss: 0.0532 - val_accuracy: 0.9895 - val_loss: 0.0375
Epoch 3/5
844/844          29s 35ms/step -
accuracy: 0.9890 - loss: 0.0374 - val_accuracy: 0.9903 - val_loss: 0.0374
Epoch 4/5
844/844          29s 35ms/step -
accuracy: 0.9910 - loss: 0.0281 - val_accuracy: 0.9903 - val_loss: 0.0393
Epoch 5/5
844/844          28s 33ms/step -

```

accuracy: 0.9933 - loss: 0.0219 - val_accuracy: 0.9905 - val_loss: 0.0370
313/313 - 3s - 8ms/step - accuracy: 0.9901 - loss: 0.0292
Test accuracy: 0.9901





```
[6]: import numpy as np
import matplotlib.pyplot as plt

def load_mnist_images(file_path):
    with open(file_path, 'rb') as f:
        # Skip the magic number and dimensions
        f.read(16)
        # Read the image data
        images = np.frombuffer(f.read(), dtype=np.uint8)
        images = images.reshape(-1, 28, 28, 1) # Reshape to 28x28x1
    return images

def load_mnist_labels(file_path):
    with open(file_path, 'rb') as f:

        f.read(8)
        # Read the label data
        labels = np.frombuffer(f.read(), dtype=np.uint8)
    return labels
```

```

# Replace these paths with the correct paths to your extracted dataset files
train_images_path = 'train-images.idx3-ubyte'
train_labels_path = 'train-labels.idx1-ubyte'

# Load the data
train_images = load_mnist_images(train_images_path)
train_labels = load_mnist_labels(train_labels_path)

# Normalize the images to the range [0, 1]
train_images = train_images.astype('float32') / 255

# Display a few samples from the dataset
fig, axes = plt.subplots(3, 5, figsize=(10, 6))
axes = axes.flatten()


# Sample indices to visualize
sample_indices = np.random.choice(range(len(train_images)), 15, replace=False)

for i, idx in enumerate(sample_indices):
    image = train_images[idx].reshape(28, 28)
    label = train_labels[idx] # Get the label as an integer

    axes[i].imshow(image, cmap='gray')
    axes[i].set_title(f'Label: {label}')
    axes[i].axis('off')

plt.tight_layout()
plt.show()

```



```
[13]: import os
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.optimizers import Adam
import matplotlib.pyplot as plt

dataset_dir = 'C:/Users/HP/Documents/coin_dataset'
batch_size = 32
img_height, img_width = 224, 224
epochs = 5

train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    validation_split=0.2)
```

```

train_generator = train_datagen.flow_from_directory(
    dataset_dir,
    target_size=(img_height, img_width),
    batch_size=batch_size,
    class_mode='categorical',
    subset='training')

validation_generator = train_datagen.flow_from_directory(
    dataset_dir,
    target_size=(img_height, img_width),
    batch_size=batch_size,
    class_mode='categorical',
    subset='validation')

num_classes = len(train_generator.class_indices)
print(f"Detected number of classes: {num_classes}")

model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(img_height, img_width, 3)),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(128, (3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dense(512, activation='relu'),
    Dropout(0.5),
    Dense(num_classes, activation='softmax') # Output layer with dynamic
    num_classes
])

model.compile(optimizer=Adam(), loss='categorical_crossentropy',
    metrics=['accuracy'])

history = model.fit(
    train_generator,
    epochs=epochs,
    validation_data=validation_generator)

model.save('inr_coin_model.h5')

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

```

```

loss = history.history['loss']
val_loss = history.history['val_loss']
epochs_range = range(epochs)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()

```

Found 720 images belonging to 1 classes.

Found 180 images belonging to 1 classes.

Detected number of classes: 1

C:\Users\HP\anaconda3\Lib\site-

packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Epoch 1/5

23/23 136s 5s/step -

accuracy: 1.0000 - loss: 0.0000e+00 - val_accuracy: 1.0000 - val_loss: 0.0000e+00

Epoch 2/5

23/23 126s 4s/step -

accuracy: 1.0000 - loss: 0.0000e+00 - val_accuracy: 1.0000 - val_loss: 0.0000e+00

Epoch 3/5

23/23 125s 4s/step -

accuracy: 1.0000 - loss: 0.0000e+00 - val_accuracy: 1.0000 - val_loss: 0.0000e+00

Epoch 4/5

23/23 123s 4s/step -

accuracy: 1.0000 - loss: 0.0000e+00 - val_accuracy: 1.0000 - val_loss: 0.0000e+00

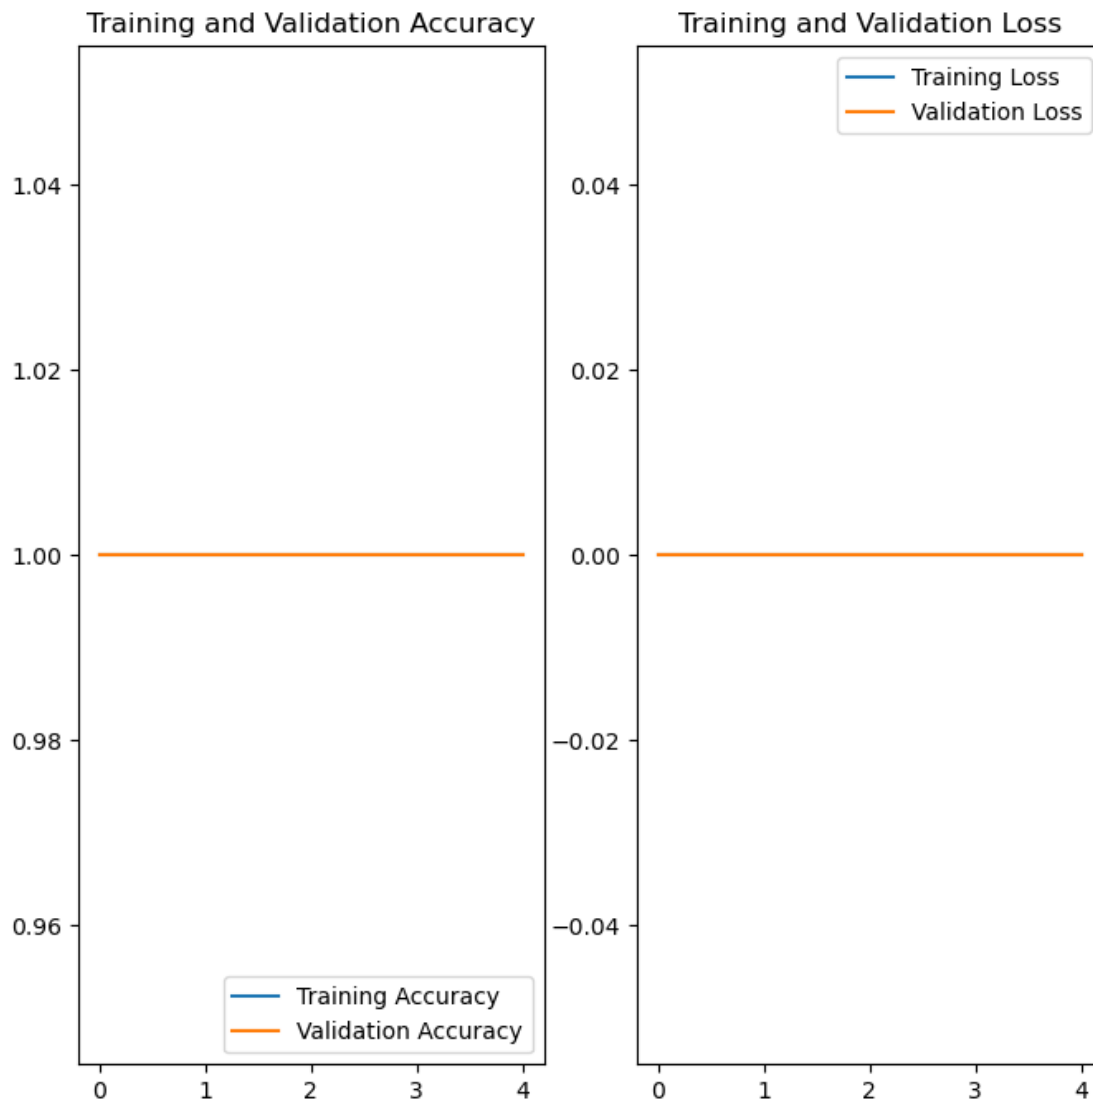
Epoch 5/5

23/23 122s 4s/step -

accuracy: 1.0000 - loss: 0.0000e+00 - val_accuracy: 1.0000 - val_loss:

0.0000e+00

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.



```
[16]: import os
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential, load_model
```

```

from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, \
↳Dropout
from tensorflow.keras.optimizers import Adam
import matplotlib.pyplot as plt
import cv2
import numpy as np
from tkinter import Tk, filedialog
from PIL import Image

dataset_dir = 'D:/AI&ML/DataSet'
model_path = 'inr_coin_model.h5'
batch_size = 32
img_height, img_width = 224, 224
epochs = 10

if not os.path.exists(model_path):
    print("Model not found, training a new model...")

    train_datagen = ImageDataGenerator(
        rescale=1./255,
        rotation_range=20,
        width_shift_range=0.2,
        height_shift_range=0.2,
        shear_range=0.2,
        zoom_range=0.2,
        horizontal_flip=True,
        validation_split=0.2)

    train_generator = train_datagen.flow_from_directory(
        dataset_dir,
        target_size=(img_height, img_width),
        batch_size=batch_size,
        class_mode='categorical',
        subset='training')

    validation_generator = train_datagen.flow_from_directory(
        dataset_dir,
        target_size=(img_height, img_width),
        batch_size=batch_size,
        class_mode='categorical',
        subset='validation')

    num_classes = len(train_generator.class_indices)
    print(f"Detected number of classes: {num_classes}")

```

```

model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(img_height,
↪img_width, 3)),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(128, (3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dense(512, activation='relu'),
    Dropout(0.5),
    Dense(num_classes, activation='softmax')
])

model.compile(optimizer=Adam(), loss='categorical_crossentropy',
↪metrics=['accuracy'])

history = model.fit(
    train_generator,
    epochs=epochs,
    validation_data=validation_generator)

model.save(model_path)

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs_range = range(epochs)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()

else:
    print("Model found, loading the model...")

```

```

model = load_model(model_path)
class_labels = ['1 INR', '2 INR', '5 INR', '10 INR']

def preprocess_image(image_path):
    img = cv2.imread(image_path)
    img = cv2.resize(img, (224, 224))
    img = img.astype('float32') / 255.0
    img = np.expand_dims(img, axis=0)
    return img

def predict_coin(image_path):
    processed_image = preprocess_image(image_path)
    predictions = model.predict(processed_image)
    predicted_class = np.argmax(predictions, axis=1)[0]
    return class_labels[predicted_class]

def upload_image():
    root = Tk()
    root.withdraw()
    image_path = filedialog.askopenfilename(title="Select an Image",
    ↪filetypes=[("Image Files", ".jpg;.jpeg;*.png")])
    if image_path:
        coin_type = predict_coin(image_path)
        print(f'The coin is: {coin_type}')

        img = cv2.imread(image_path)
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
        plt.imshow(img)
        plt.title(f"Coin Type: {coin_type}")
        plt.axis('off')
        plt.show()

upload_image()

```

Model found, loading the model...

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

1/1 0s 118ms/step

The coin is: 1 INR

Coin Type: 1 INR



```
[1]: import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.callbacks import ModelCheckpoint

(train_images, train_labels), (test_images, test_labels) = mnist.load_data()

train_images = train_images.reshape((train_images.shape[0], 28, 28, 1)).
    ↪astype('float32') / 255
test_images = test_images.reshape((test_images.shape[0], 28, 28, 1)).
    ↪astype('float32') / 255

train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)

model = Sequential()
model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(MaxPooling2D((2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu'))
```

```

model.add(MaxPooling2D((2, 2)))
model.add(Flatten())
model.add(Dense(64, activation='relu'))
model.add(Dense(10, activation='softmax'))

model.compile(optimizer='adam', loss='categorical_crossentropy',
              metrics=['accuracy'])

checkpoint = ModelCheckpoint('mnist_cnn_model.keras', save_best_only=True)

history = model.fit(train_images, train_labels, epochs=5, batch_size=64,
                  validation_data=(test_images, test_labels), callbacks=[checkpoint])

test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f'Test accuracy: {test_acc:.4f}')

```

C:\Users\HP\anaconda3\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Epoch 1/5

938/938 14s 14ms/step -

accuracy: 0.8831 - loss: 0.3961 - val_accuracy: 0.9817 - val_loss: 0.0560

Epoch 2/5

938/938 14s 15ms/step -

accuracy: 0.9833 - loss: 0.0546 - val_accuracy: 0.9872 - val_loss: 0.0383

Epoch 3/5

938/938 13s 13ms/step -

accuracy: 0.9883 - loss: 0.0366 - val_accuracy: 0.9871 - val_loss: 0.0418

Epoch 4/5

938/938 13s 13ms/step -

accuracy: 0.9913 - loss: 0.0284 - val_accuracy: 0.9859 - val_loss: 0.0482

Epoch 5/5

938/938 13s 13ms/step -

accuracy: 0.9929 - loss: 0.0214 - val_accuracy: 0.9901 - val_loss: 0.0297

313/313 1s 3ms/step -

accuracy: 0.9868 - loss: 0.0400

Test accuracy: 0.9901

```

[4]: import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.models import load_model
from PIL import Image

model = load_model('mnist_cnn_model.keras')

```

```

def load_and_preprocess_image(img_path):
    img = Image.open(img_path).convert('L')
    img = img.resize((28, 28))
    img_array = np.array(img).astype('float32') / 255
    img_array = np.expand_dims(img_array, axis=0)
    img_array = np.expand_dims(img_array, axis=-1)
    return img_array

img_path = 'C:/Users/HP/Downloads/number7.png'
processed_image = load_and_preprocess_image(img_path)

prediction = model.predict(processed_image)
predicted_label = np.argmax(prediction)

print(f'The predicted label is: {predicted_label}')

plt.imshow(processed_image[0].squeeze(), cmap='gray')
plt.title(f'Predicted: {predicted_label}')
plt.axis('off')
plt.show()

```

1/1 0s 63ms/step
The predicted label is: 7

Predicted: 7



[]: