

week6

September 25, 2024

[1]:

```
[1]: import pandas as pd
credit_df = pd.read_csv("C:/Users/HP/Downloads/credit.csv")
credit_df
```

```
[1]:      id  Income  Limit  Rating  Cards  Age  Education  Gender  Student  \
0      1   14.891   3606    283      2   34          11    Male      No
1      2  106.025   6645    483      3   82          15  Female      Yes
2      3  104.593   7075    514      4   71          11    Male      No
3      4  148.924   9504    681      3   36          11  Female      No
4      5   55.882   4897    357      2   68          16    Male      No
..  ...  ...  ...  ...  ...  ...  ...  ...  ...
395  396   12.096   4100    307      3   32          13    Male      No
396  397   13.364   3838    296      5   65          17    Male      No
397  398   57.872   4171    321      5   67          12  Female      No
398  399   37.728   2525    192      1   44          13    Male      No
399  400   18.701   5524    415      5   64           7  Female      No
```

	Married	Ethnicity	Balance
0	Yes	Caucasian	333
1	Yes	Asian	903
2	No	Asian	580
3	No	Asian	964
4	Yes	Caucasian	331
..
395	Yes	Caucasian	560
396	No	African American	480
397	Yes	Caucasian	138
398	Yes	Caucasian	0
399	No	Asian	966

[400 rows x 12 columns]

```
[6]: credit_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
```


Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	id	400 non-null	int64
1	Income	400 non-null	float64
2	Limit	400 non-null	int64
3	Rating	400 non-null	int64
4	Cards	400 non-null	int64
5	Age	400 non-null	int64
6	Education	400 non-null	int64
7	Gender	400 non-null	object
8	Student	400 non-null	object
9	Married	400 non-null	object
10	Ethnicity	400 non-null	object
11	Balance	400 non-null	int64

dtypes: float64(1), int64(7), object(4)
memory usage: 37.6+ KB


```
[5]: credit_df['Rating'].plot(kind='box')
```

```
[5]: <Axes: >
```



```
[7]: credit_df['Rating'].plot(kind='hist')
```


```
[7]: <Axes: ylabel='Frequency'>
```



```
[8]: import numpy as np
import matplotlib.pyplot as plt
np.random.seed(102)
grades = np.concatenate([[50,52,53,55,56,60,61,62,65,67]*20,
np.random.randint(0, 101, size=300)])
Q1 = np.percentile(grades , 25)
Q3 = np.percentile(grades , 75)
Q1,Q3 = np.percentile(grades , [25,75])
IQR = Q3 - Q1
ul = Q3+1.5*IQR
ll = Q1-1.5*IQR
outliers = grades[(grades > ul) | (grades < ll)]
print(outliers)
fig = plt.figure(figsize=(6,5))
hypo = np.random.randint(20, 81, size=500)
plt.boxplot([grades, hypo], widths=0.5)
plt.xticks([1,2],['Our data', 'Hypothetical data'])
plt.ylabel('Grade')
```

```
plt.title('Box plot of midterm grade')
plt.show()
```

```
[ 0  7  4  3  0  4  2  7  6 100  1  3  0  3 100 100 100 100
  4  0  3  6  6  6 100  7  6 100 100  6  3  6  1  6  0]
```



```
[11]: import numpy as np

data = [1, 2, 2, 2, 3, 1, 1, 15, 2, 2, 2, 3, 1, 1, 2]
mean = np.mean(data)
std = np.std(data)

print('Mean of the dataset is:', mean)
print('Standard deviation is:', std)

threshold = 3
outliers = []

for i in data:
    z = (i - mean) / std
```

```

    if abs(z) > threshold:
        outliers.append(i)

print('Outliers in dataset based on Z-score are:', outliers)

```

Mean of the dataset is: 2.6666666666666665
Standard deviation is: 3.3598941782277745
Outliers in dataset based on Z-score are: [15]

```

[12]: q1 = credit_df["Age"].quantile(0.25)
      q3 = credit_df['Age'].quantile(0.75)
      iqr = q3-q1
      upper_bound = q3+(1.5*iqr)
      lower_bound = q1-(1.5*iqr)

```

```

[14]: upperIndex = credit_df[credit_df['Age']>upper_bound].index
      credit_df.drop(upperIndex,inplace=True)
      lowerIndex = credit_df[credit_df['Age']<lower_bound].index
      credit_df.drop(lowerIndex,inplace=True)
      credit_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Unnamed: 0      400 non-null   int64
1   Income          400 non-null   float64
2   Limit           400 non-null   int64
3   Rating          400 non-null   int64
4   Cards           400 non-null   int64
5   Age             400 non-null   int64
6   Education       400 non-null   int64
7   Gender          400 non-null   object
8   Student         400 non-null   object
9   Married         400 non-null   object
10  Ethnicity       400 non-null   object
11  Balance         400 non-null   int64
dtypes: float64(1), int64(7), object(4)
memory usage: 37.6+ KB

```

```

[16]: m = np.mean(credit_df['Age'])
      print('mean:',m)
      for i in credit_df['Age']:
          if i<lower_bound or i>upper_bound :
              titanic_df['Age'] = titanic_df['Age'].replace(i,m)

```

mean: 55.6675

```
[17]: m = credit_df['Age'].median()
print("median",m)
for i in credit_df['Age']:
    if i<lower_bound or i>upper_bound :
        credit_df['Age'] = credit_df['Age'].replace(i,m)
```

median 56.0

```
[18]: for i in credit_df['Age']:
    if i<lower_bound or i>upper_bound :
        credit_df['Age'] = credit_df['Age'].replace(i,0)
```

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import math
```

```
[5]: card_approval_df=pd.read_csv("C:/Users/HP/Downloads/credit.csv")
print(card_approval_df.head())
```

	Unnamed: 0	Income	Limit	Rating	Cards	Age	Education	Gender	Student	\
0	1	14.891	3606	283	2	34	11	Male	No	
1	2	106.025	6645	483	3	82	15	Female	Yes	
2	3	104.593	7075	514	4	71	11	Male	No	
3	4	148.924	9504	681	3	36	11	Female	No	
4	5	55.882	4897	357	2	68	16	Male	No	

	Married	Ethnicity	Balance
0	Yes	Caucasian	333
1	Yes	Asian	903
2	No	Asian	580
3	No	Asian	964
4	Yes	Caucasian	331

```
[6]: print(card_approval_df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Unnamed: 0      400 non-null   int64
1   Income          400 non-null   float64
2   Limit           400 non-null   int64
3   Rating          400 non-null   int64
4   Cards           400 non-null   int64
5   Age             400 non-null   int64
```

```

6   Education    400 non-null    int64
7   Gender       400 non-null    object
8   Student      400 non-null    object
9   Married      400 non-null    object
10  Ethnicity    400 non-null    object
11  Balance      400 non-null    int64
dtypes: float64(1), int64(7), object(4)
memory usage: 37.6+ KB
None

```

[7]:

[7]:

```

import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
get_ipython().run_line_magic('matplotlib', 'inline')

```

[8]:

```

boston_dataset = pd.read_csv('C:/Users/HP/Downloads/BostonHousing.csv')
boston_dataset.keys()

```

[8]:

```

Index(['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax',
       'ptratio', 'b', 'lstat', 'medv'],
      dtype='object')

```

[9]:

```

boston_dataset.head(5)

```

[9]:

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	

	b	lstat	medv
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2

[10]:

```

boston_dataset.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   crim        506 non-null    float64

```

```
1   zn      506 non-null   float64
2   indus   506 non-null   float64
3   chas    506 non-null   int64
4   nox     506 non-null   float64
5   rm      506 non-null   float64
6   age     506 non-null   float64
7   dis     506 non-null   float64
8   rad     506 non-null   int64
9   tax     506 non-null   int64
10  ptratio  506 non-null   float64
11  b       506 non-null   float64
12  lstat   506 non-null   float64
13  medv    506 non-null   float64
dtypes: float64(11), int64(3)
memory usage: 55.5 KB
```

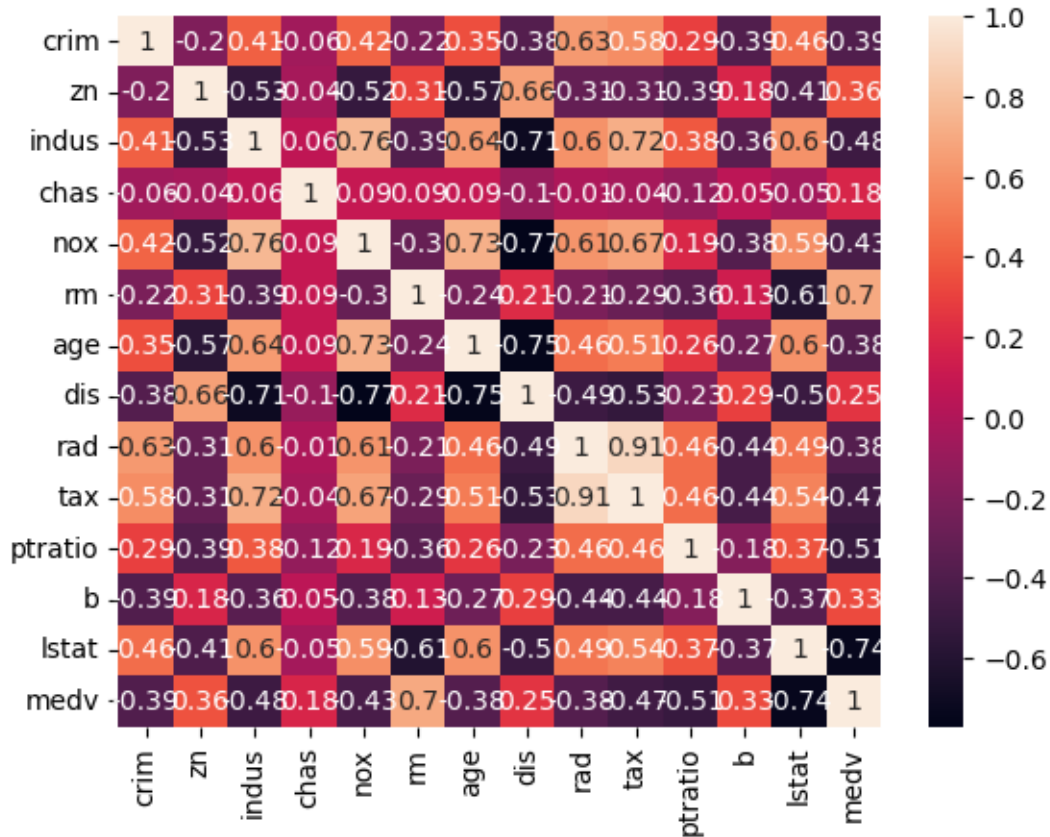
```
[11]: boston_dataset.isnull().sum()
```

```
[11]: crim      0
      zn       0
      indus    0
      chas     0
      nox      0
      rm       0
      age      0
      dis      0
      rad      0
      tax      0
      ptratio  0
      b        0
      lstat    0
      medv     0
      dtype: int64
```


```
[12]: correlation_matrix = boston_dataset.corr().round(2)
```

```
[13]: sns.heatmap(data=correlation_matrix, annot=True)
```

```
[13]: <Axes: >
```

```
[55]: import matplotlib.pyplot as plt
plt.figure(figsize=(20, 5))
features = ['lstat', 'rm']
target = boston_dataset['tax']
for i, col in enumerate(features):
    plt.subplot(1, len(features), i+1)
    x = boston_dataset[col]
    y = target
    plt.scatter(x, y, marker='o')
    plt.title(col)
    plt.xlabel(col)
    plt.ylabel('tax')
```



```
[16]: X = pd.DataFrame(np.c_[boston_dataset['tax'], boston_dataset['rm']],
    ↪columns=['lstat', 'rm'])
Y = boston_dataset['age']
```

```
[17]: from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
    ↪random_state=42)
print(X_train.shape)
print(X_test.shape)
print(Y_train.shape)
print(Y_test.shape)
```

```
(404, 2)
(102, 2)
(404,)
(102,)
```

```
[18]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
[19]: lin_model = LinearRegression()
lin_model.fit(X_train, Y_train)
```

```
[19]: LinearRegression()
```

```
[21]: y_train_predict = lin_model.predict(X_train)
rmse = (np.sqrt(mean_squared_error(Y_train, y_train_predict)))
r2 = r2_score(Y_train, y_train_predict)
print("The model performance for training set")
print("-----")
print('RMSE is {}'.format(rmse))
print('R2 score is {}'.format(r2))
print("\n")
# model evaluation for testing set
y_test_predict = lin_model.predict(X_test)
```

```

# root mean square error of the model
rmse = (np.sqrt(mean_squared_error(Y_test, y_test_predict)))
# r-squared score of the model
r2 = r2_score(Y_test, y_test_predict)
print("The model performance for testing set")
print("-----")
print('RMSE is {}'.format(rmse))
print('R2 score is {}'.format(r2))

```

The model performance for training set

RMSE is 24.54898757784132

R2 score is 0.2291232217092818

The model performance for testing set


RMSE is 22.33059428323184

R2 score is 0.39666537877593555


```

[22]: plt.scatter(Y_test, y_test_predict)
plt.show()

```



```
[24]: plt.figure(figsize=(20, 5))
features = ['lstat', 'rm']
target = boston_dataset['age']
for i, col in enumerate(features):
    plt.subplot(1, len(features), i+1)
    x = boston_dataset[col]
    y = target
    plt.scatter(x, y, marker='o')
    plt.title(col)
    plt.xlabel(col)
    plt.ylabel('age')
```




```
[ ]: import numpy as np
import matplotlib.pyplot as plt
```

```
[28]: %matplotlib inline
def gradient_descent(x,y):
    m = b = 1
    rate = 0.01
    n = len(x)
    plt.scatter(x,y)
    for i in range(100):
        y_predicted = m * x + b
        plt.plot(x,y_predicted,color='green')
        md = -(2/n)*sum(x*(y-y_predicted))
        yd = -(2/n)*sum(y-y_predicted)
        m = m - rate * md
        b = b - rate * yd
```

```
[29]: x = np.array([1,2,3,4,5])
y = np.array([5,7,9,11,13])
```

```
[30]: gradient_descent(x,y)
```



```
[31]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
# Importing the dataset
datas = pd.read_csv('C:/Users/HP/Downloads/BostonHousing.csv')
datas
```

```
[31]:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	
..	
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	

	ptratio	b	lstat	medv
0	15.3	396.90	4.98	24.0

1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2
..
501	21.0	391.99	9.67	22.4
502	21.0	396.90	9.08	20.6
503	21.0	396.90	5.64	23.9
504	21.0	393.45	6.48	22.0
505	21.0	396.90	7.88	11.9

[506 rows x 14 columns]

```
[33]: X = datas.iloc[:, 1:2].values
      y = datas.iloc[:, 2].values
```


```
[34]: from sklearn.linear_model import LinearRegression
      lin = LinearRegression()
      lin.fit(X, y)
```

[34]: LinearRegression()


```
[35]: from sklearn.preprocessing import PolynomialFeatures
      poly = PolynomialFeatures(degree = 4)+
      X_poly = poly.fit_transform(X)
      poly.fit(X_poly, y)
      lin2 = LinearRegression()
      lin2.fit(X_poly, y)
```

[35]: LinearRegression()

```
[36]: plt.scatter(X, y, color = 'blue')
      plt.plot(X, lin.predict(X), color = 'red')
      plt.title('Linear Regression')
      plt.xlabel('Temperature')
      plt.ylabel('Pressure')
      plt.show()
```



```
[37]: plt.scatter(X, y, color = 'blue')
plt.plot(X, lin2.predict(poly.fit_transform(X)), color = 'red')
plt.title('Polynomial Regression')
plt.xlabel('Temperature')
plt.ylabel('Pressure')
plt.show()
```



```
[38]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
[39]: from sklearn.model_selection import train_test_split
```

```
[42]: df = pd.read_csv('C:/Users/HP/Downloads/Advertising.csv.')
df
```

```
[42]:
```

	ID	TV	Radio	Newspaper	Sales
0	1	230.1	37.8	69.2	22.1
1	2	44.5	39.3	45.1	10.4
2	3	17.2	45.9	69.3	9.3
3	4	151.5	41.3	58.5	18.5
4	5	180.8	10.8	58.4	12.9
..
195	196	38.2	3.7	13.8	7.6
196	197	94.2	4.9	8.1	9.7


```

197 198 177.0    9.3        6.4  12.8
198 199 283.6   42.0       66.2  25.5
199 200 232.1    8.6        8.7  13.4

```

[200 rows x 5 columns]

```
[43]: df.dropna(inplace=True,axis=0)
df
```

```
[43]:
```

	ID	TV	Radio	Newspaper	Sales
0	1	230.1	37.8	69.2	22.1
1	2	44.5	39.3	45.1	10.4
2	3	17.2	45.9	69.3	9.3
3	4	151.5	41.3	58.5	18.5
4	5	180.8	10.8	58.4	12.9
..
195	196	38.2	3.7	13.8	7.6
196	197	94.2	4.9	8.1	9.7
197	198	177.0	9.3	6.4	12.8
198	199	283.6	42.0	66.2	25.5
199	200	232.1	8.6	8.7	13.4

[200 rows x 5 columns]

```
[45]: y = df['TV']
X = df.drop('TV',axis=1)
```

```
[46]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.
↪3,random_state=101)
```

```
[47]: scaler = StandardScaler()
scaler.fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

```
[48]: lr = LinearRegression()
model = lr.fit(X_train,y_train)
```

```
[49]: y_pred = model.predict(X_test)
ydf = pd.DataFrame({'y_test':y_test,'y_pred':y_pred})
rslt_df = ydf.sort_values(by = 'y_test')
```

```
[50]: print(mean_squared_error(y_test,y_pred))
```


933.1209747971208

```
[51]: print(r2_score(y_test, y_pred))
```

0.8790644106224436

```
[52]: import matplotlib.pyplot as plt
plt.scatter(ydf['y_test'], ydf['y_pred'])
```

```
[52]: <matplotlib.collections.PathCollection at 0x20664920ad0>
```



```
[53]: model.coef_
```

```
[53]: array([ 4.19811629, -49.89190342,  2.02818532,  94.1700075 ])
```

```
[54]: model.intercept_
```

```
[54]: 151.66071428571433
```

```
[7]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Load the dataset from a file
# Replace 'your_dataset.csv' with the actual path to your file
data = pd.read_csv("C:/Users/HP/Downloads/credit.csv")
```

```

bus = pd.read_csv("C:/Users/HP/Downloads/marksheet.csv")
# Assuming the dataset has columns 'feature1', 'feature2', ..., 'target'
# Replace these with the actual column names
X = data[['Education', 'Balance', 'Rating']] # Feature columns
Y = data['Income'] # Target column

# Split the dataset into training and testing sets
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
    random_state=42)

# Train the Linear Regression model
lin_model = LinearRegression()
lin_model.fit(X_train, Y_train)

# Model evaluation for training set
y_train_predict = lin_model.predict(X_train)
rmse_train = np.sqrt(mean_squared_error(Y_train, y_train_predict))
r2_train = r2_score(Y_train, y_train_predict)

print("The model performance for training set")
print("-----")
print('RMSE is {}'.format(rmse_train))
print('R2 score is {}'.format(r2_train))
print("\n")

# Model evaluation for testing set
y_test_predict = lin_model.predict(X_test)
rmse_test = np.sqrt(mean_squared_error(Y_test, y_test_predict))
r2_test = r2_score(Y_test, y_test_predict)

print("The model performance for testing set")
print("-----")
print('RMSE is {}'.format(rmse_test))
print('R2 score is {}'.format(r2_test))

```

The model performance for training set

RMSE is 15.02117077384028

R2 score is 0.8185903172179333

The model performance for testing set

RMSE is 15.36793364811597

R2 score is 0.8064320939429278

```
[12]: credit_df = pd.read_csv("C:/Users/HP/Downloads/credit.csv")
credit_df
```

```
[12]:      Unnamed: 0   Income  Limit  Rating  Cards  Age  Education  Gender  \
0             1   14.891   3606    283     2   34           11    Male
1             2  106.025   6645    483     3   82           15  Female
2             3  104.593   7075    514     4   71           11    Male
3             4  148.924   9504    681     3   36           11  Female
4             5   55.882   4897    357     2   68           16    Male
..          ...     ...     ...     ...     ...   ...     ...
395          396   12.096   4100    307     3   32           13    Male
396          397   13.364   3838    296     5   65           17    Male
397          398   57.872   4171    321     5   67           12  Female
398          399   37.728   2525    192     1   44           13    Male
399          400   18.701   5524    415     5   64            7  Female
```

```
      Student Married      Ethnicity  Balance
0         No      Yes      Caucasian    333
1         Yes     Yes          Asian    903
2         No     No          Asian    580
3         No     No          Asian    964
4         No     Yes      Caucasian    331
..          ...     ...           ...     ...
395        No     Yes      Caucasian    560
396        No     No  African American    480
397        No     Yes      Caucasian    138
398        No     Yes      Caucasian     0
399        No     No          Asian    966
```

[400 rows x 12 columns]

```
[13]: credit_df = pd.read_csv("C:/Users/HP/Downloads/marksheet.csv")
credit_df
```

```
[13]:      id   Name  Gender  Age  Section  Science  English  History  Maths
0      1  Bronnie  Female   13      C        21       81       62     49
1      2   Lemmie   Male   15      B        29       41       17     40
2      3    Danya  Female   14      C        12       87       16     96
3      4    Denna  Female   14      B        15       53       82     33
4      5  Jocelin   Male   14      A        43        6        3     21
..   ...   ...     ...   ...     ...     ...     ...     ...
245  246   Nickie   Male   13      C        28       15       25     10
246  247     Rog  Female   13      B         1        4       68     65
247  248    Kaia   Male   15      B       93       48       82     44
248  249    Anni  Female   14      B       35       73       66     59
249  250  Fernande   Male   15      B       50        8       80     78
```

[250 rows x 9 columns]

```
[19]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Load the datasets
data_credit = pd.read_csv("C:/Users/HP/Downloads/credit.csv")
data_marksheet = pd.read_csv("C:/Users/HP/Downloads/marksheet.csv")

# Display the first few rows of each dataset to understand their structure
print("Credit Dataset:")
print(data_credit.head())
print("\nMarksheet Dataset:")
print(data_marksheet.head())

# For the sake of this example, let's assume both datasets share a common
↳ column 'ID'
# and we want to predict 'Income' from the credit dataset using features from
↳ both datasets.

# Merge the datasets on a common column 'ID'
merged_data = pd.merge(data_credit, data_marksheet, on='id')

# Select features from both datasets
X = merged_data[['Education', 'Balance', 'Rating', 'Maths', 'English']] #
↳ Example features
Y = merged_data['Income'] # Target variable

# Split the combined dataset into training and testing sets
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
↳ random_state=42)

# Train the Linear Regression model
lin_model = LinearRegression()
lin_model.fit(X_train, Y_train)

# Model evaluation for training set
y_train_predict = lin_model.predict(X_train)
rmse_train = np.sqrt(mean_squared_error(Y_train, y_train_predict))
r2_train = r2_score(Y_train, y_train_predict)

print("\nThe model performance for the training set")
print("-----")
print('RMSE is {}'.format(rmse_train))
```

```

print('R2 score is {}'.format(r2_train))

# Model evaluation for testing set
y_test_predict = lin_model.predict(X_test)
rmse_test = np.sqrt(mean_squared_error(Y_test, y_test_predict))
r2_test = r2_score(Y_test, y_test_predict)

print("\nThe model performance for the testing set")
print("-----")
print('RMSE is {}'.format(rmse_test))
print('R2 score is {}'.format(r2_test))

```

Credit Dataset:

	id	Income	Limit	Rating	Cards	Age	Education	Gender	Student	Married	\
0	1	14.891	3606	283	2	34	11	Male	No	Yes	
1	2	106.025	6645	483	3	82	15	Female	Yes	Yes	
2	3	104.593	7075	514	4	71	11	Male	No	No	
3	4	148.924	9504	681	3	36	11	Female	No	No	
4	5	55.882	4897	357	2	68	16	Male	No	Yes	

		Ethnicity	Balance
0		Caucasian	333
1		Asian	903
2		Asian	580
3		Asian	964
4		Caucasian	331

Marksheet Dataset:

	id	Name	Gender	Age	Section	Science	English	History	Maths
0	1	Bronnie	Female	13	C	21	81	62	49
1	2	Lemmie	Male	15	B	29	41	17	40
2	3	Danya	Female	14	C	12	87	16	96
3	4	Denna	Female	14	B	15	53	82	33
4	5	Jocelin	Male	14	A	43	6	3	21

The model performance for the training set

RMSE is 15.284125890822864

R2 score is 0.7910326438413708

The model performance for the testing set

RMSE is 13.685085424490035

R2 score is 0.8125264172047937

```

[15]: credit_df = pd.read_csv("C:/Users/HP/Downloads/credit.csv")
      credit_df

```

```
[15]:
```

	id	Income	Limit	Rating	Cards	Age	Education	Gender	Student	\
0	1	14.891	3606	283	2	34	11	Male	No	
1	2	106.025	6645	483	3	82	15	Female	Yes	
2	3	104.593	7075	514	4	71	11	Male	No	
3	4	148.924	9504	681	3	36	11	Female	No	
4	5	55.882	4897	357	2	68	16	Male	No	
..	
395	396	12.096	4100	307	3	32	13	Male	No	
396	397	13.364	3838	296	5	65	17	Male	No	
397	398	57.872	4171	321	5	67	12	Female	No	
398	399	37.728	2525	192	1	44	13	Male	No	
399	400	18.701	5524	415	5	64	7	Female	No	

	Married	Ethnicity	Balance
0	Yes	Caucasian	333
1	Yes	Asian	903
2	No	Asian	580
3	No	Asian	964
4	Yes	Caucasian	331
..
395	Yes	Caucasian	560
396	No	African American	480
397	Yes	Caucasian	138
398	Yes	Caucasian	0
399	No	Asian	966

[400 rows x 12 columns]

```
[ ]:
```